F1 versus false alarms across 30 dataset variants

Each point averages 30 dataset-level means; the upper-left direction is preferable.

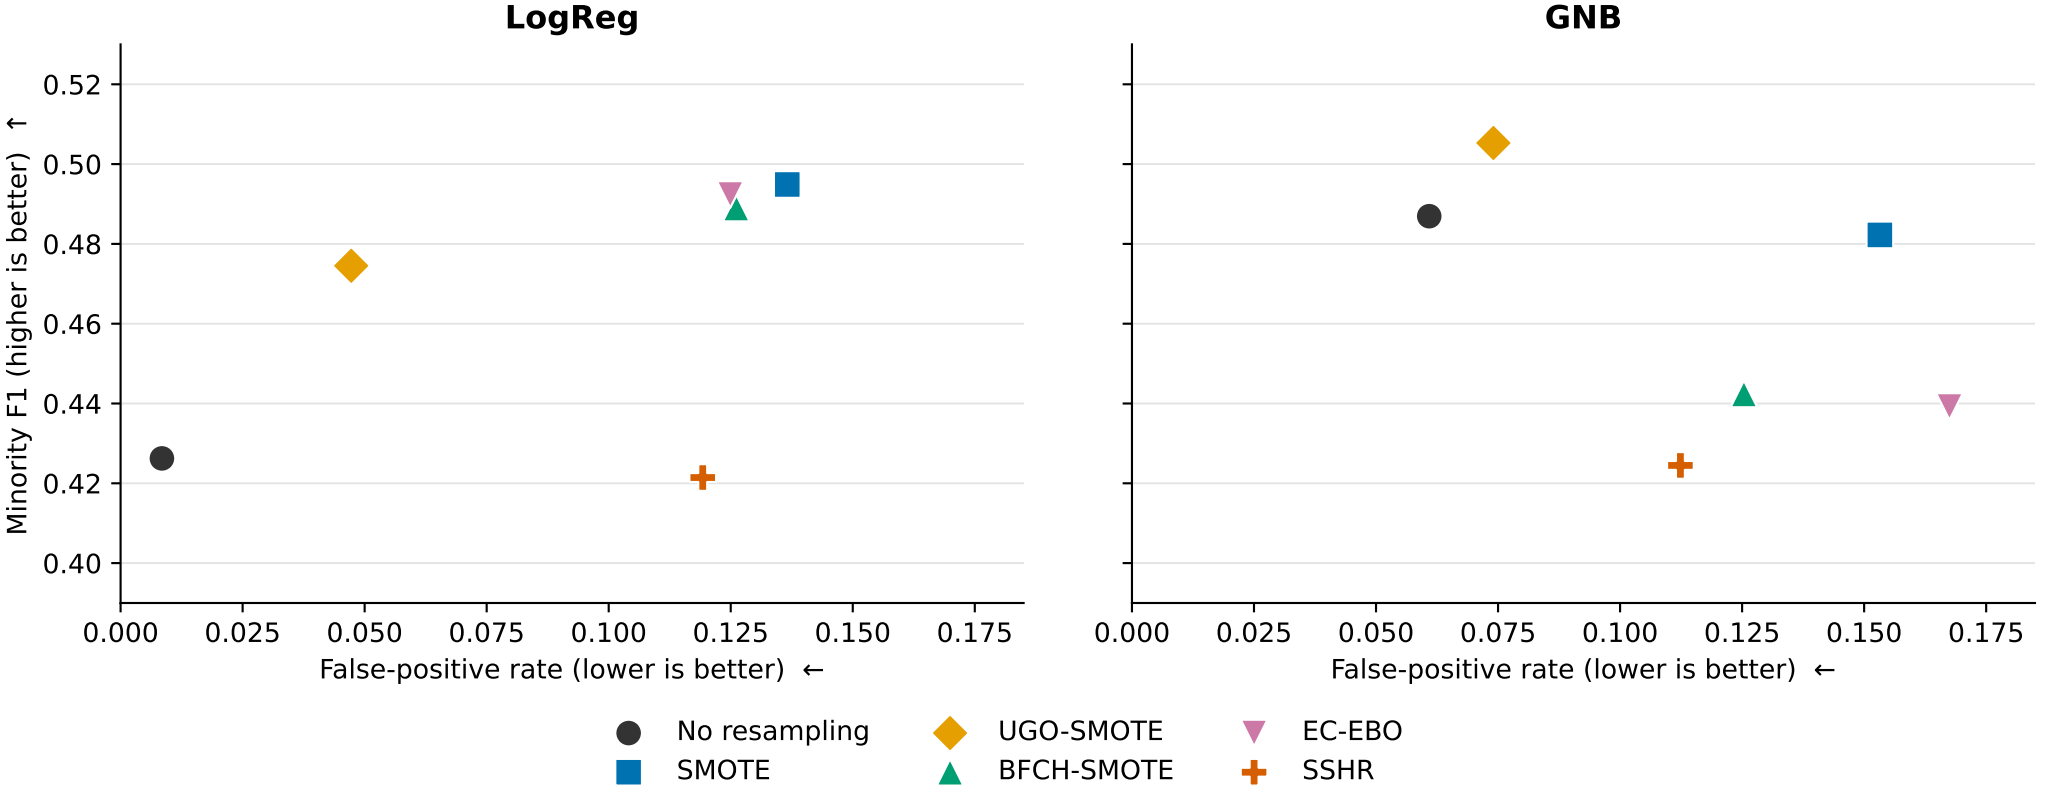

Descriptive averages; variants share source families. Other metrics remain in the report.

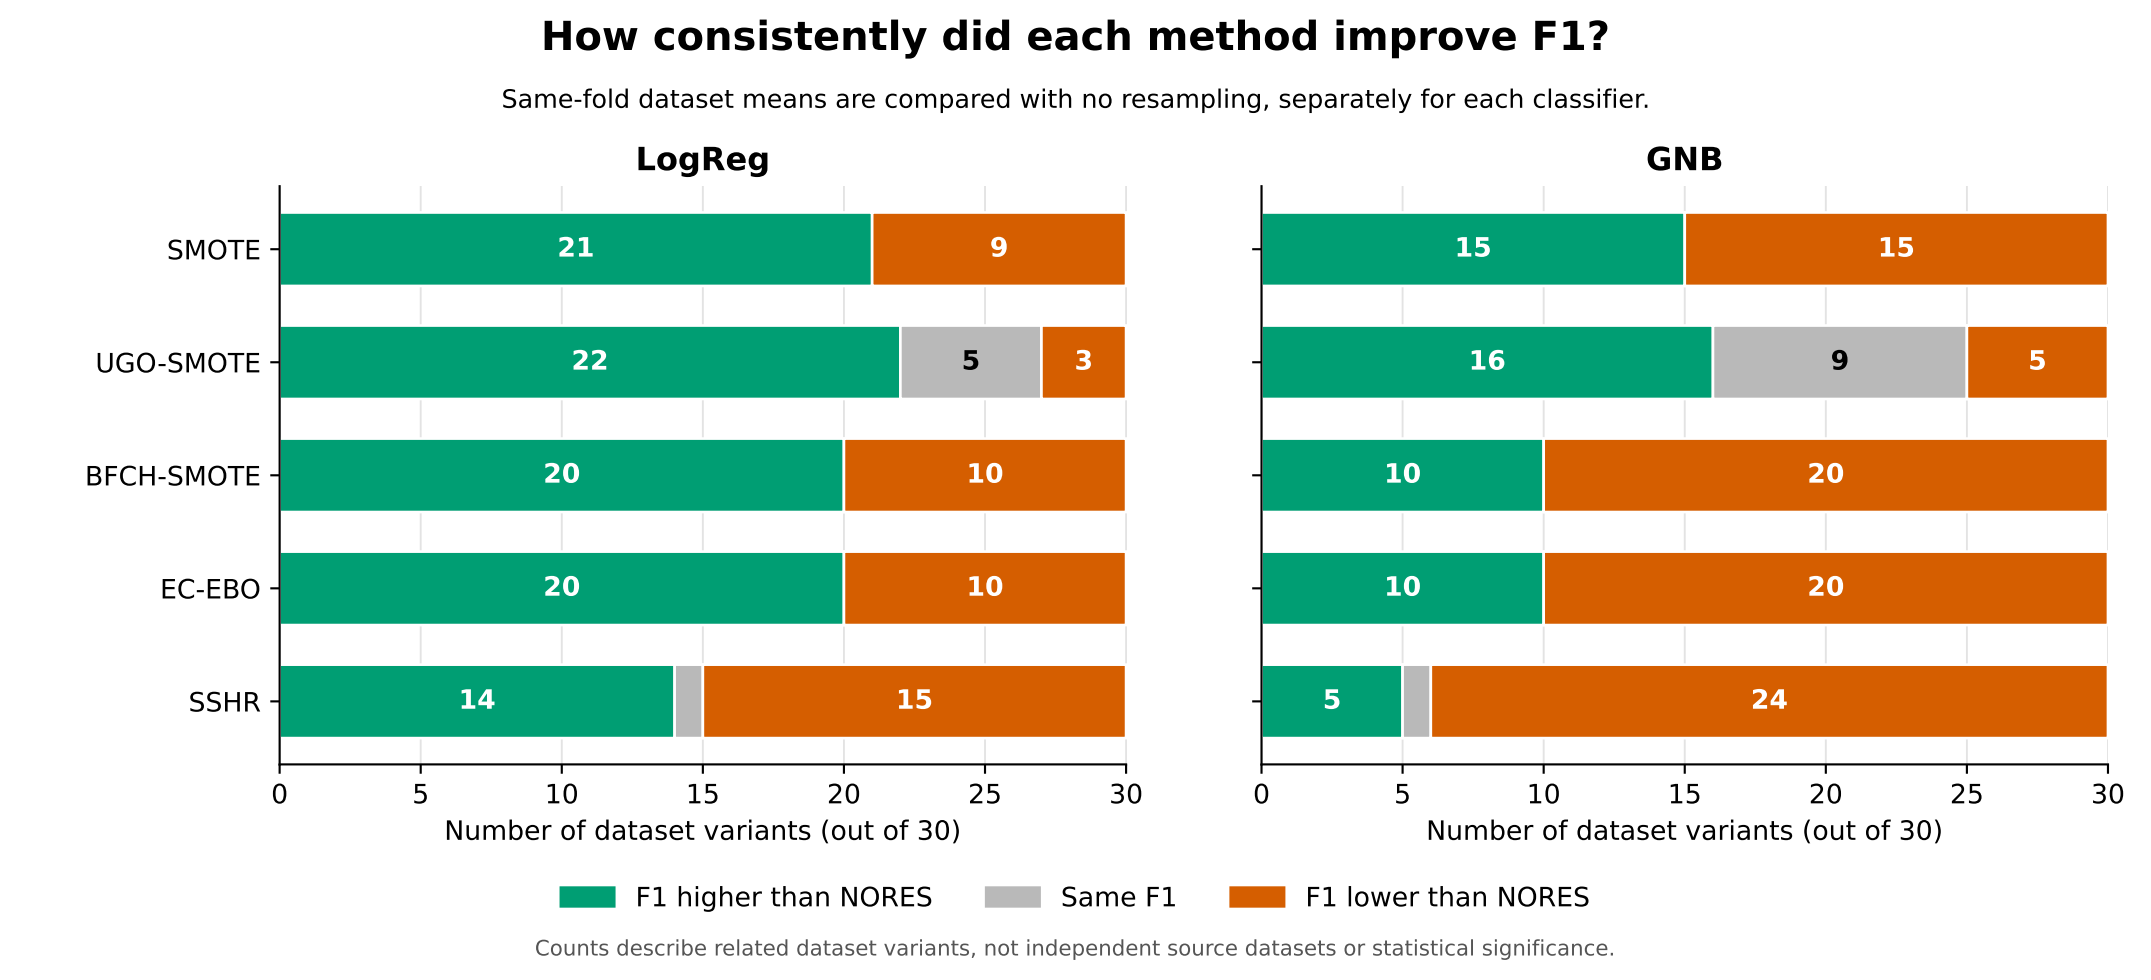
How consistently did each method improve F1?
Same-fold dataset means are compared with no resampling, separately for each classifier.
LogReg
GNB
SMOTE
UGO-SMOTE
BFCH-SMOTE
EC-EBO
SSHR
21
9
22
5
3
20
10
20
10
14
15
15
15
16
9
5
10
20
10
20
5
24
0
5
10
15
20
25
30
Number of dataset variants (out of 30)
F1 higher than NORES
Same F1
F1 lower than NORES
Counts describe related dataset variants, not independent source datasets or statistical significance.

What did the unlabelled pool change for SSHR?

Each dot is one paired dataset-variant effect; black diamonds are 30-variant means.

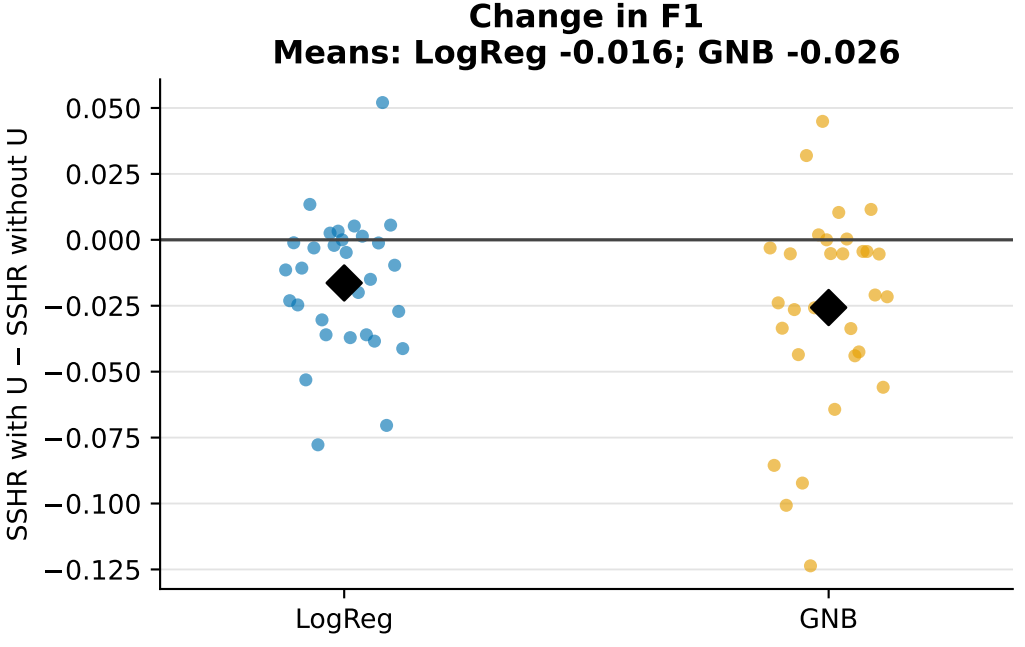

Below zero: unlabelled data lowered F1

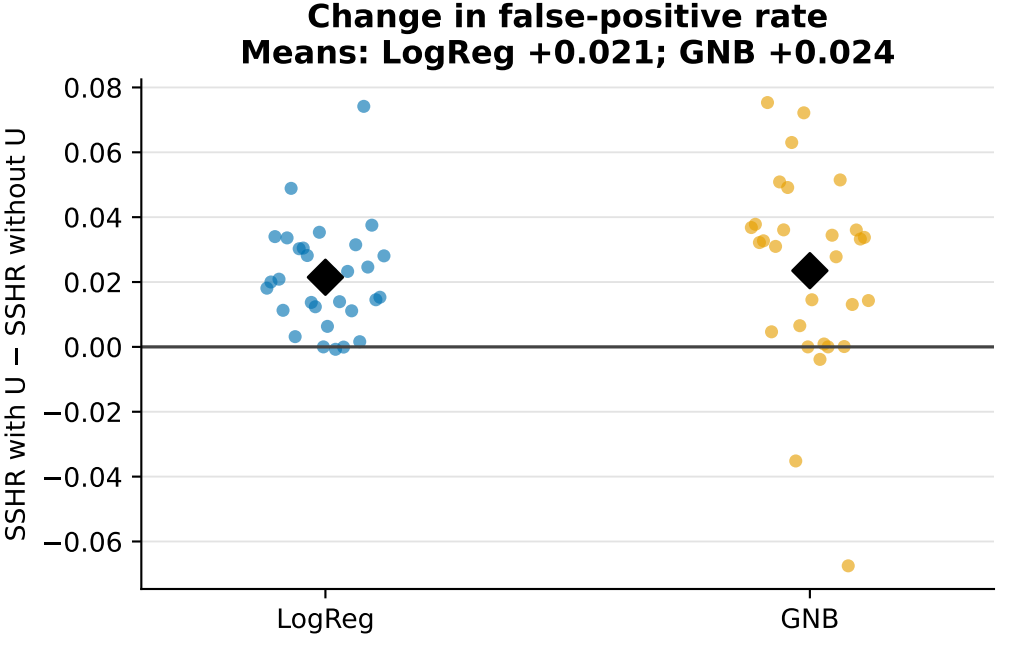

Above zero: unlabelled data raised false alarms

Same labelled pool and test folds; only access to training-only U changes. Descriptive effects.